


Writing a Go application with PostgreSQL using pgx

Henrique Vicente

Go meetup @ HATCH Amsterdam, 12 April 2022



PostgreSQL
the world's most advanced open source database



Trajectory to PostgreSQL

Migration from a document-oriented database to
PostgreSQL (RDS) + OpenSearch (Amazon's fork of Elasticsearch)

Problems with previous database

- Poor developer experience
- Reliability and performance issues
- High cost: support, licensing, and operational
- Vendor lock-in
- Steep learning curve

Why PostgreSQL

- Object-Relational Database Management System
- Almost 100% SQL standards-compliant
- Excellent for client/server architecture
- **JSON types**
- Custom data types
- Inheritance between tables

Say NoSQL

ONE MORE TIME

PostgreSQL advantages

- Great developer experience
- Reliability
- Performance
- Open-source software without vendor lock-in
- Community and multiple support providers
- Works everywhere: even on a new M1 MacBook

PostgreSQL drawbacks

- Plan schema up-front (wait, isn't this an advantage?)
- Full-text search capabilities is limited (but we've got it covered).

Environment variables for configuring PostgreSQL

<https://www.postgresql.org/docs/current/libpq-envvars.html>

```
1 export PGHOST="localhost"
2 export PGPORT=5432
3 export PGUSER="username"
4 export PGPASSWORD="your-secret-password"
5 export PGDATABASE="test_whatever"
6
7 # Set a timeout for acquiring a connection to the database:
8 export PGCONNECT_TIMEOUT=5
```


Environment variables might be problematic due to security concerns, so be careful.

- Globals
- Easy to leak to child processes or metrics probes

Also:

- Process --flags are as bad!
- Can you point to secrets, rather than expose them?

Tip: direnv makes it very convenient to use it on a development environment.

database/sql vs. pgx interface

Talking to SQL databases in Go, you almost always want to use a database through database/sql.

It provides a nice common interface, and makes supporting different databases easier.

PostgreSQL drivers for database/sql:

- github.com/jackc/pgx
- github.com/lib/pq (discontinued in favor of pgx)

database/sql vs. pgx interface

Among better performance, by using pgx interface you've:

- Support for approximately 70 different PostgreSQL types
- Automatic statement preparation and caching
- Batch queries
- Conversion of PostgreSQL arrays to Go slice mapping
- JSON and JSONB support
- COPY protocol support


database/sql is limited to int64, float64, bool, []byte, string, time.Time, or nil

Concurrency: pgxpool

- `pgx.Conn` is a low-level implementation
- Use `pgxpool.Conn` to access the database concurrently

Default:

```
conf.MaxConns = runtime.NumCPU()
```



Database migrations: 001_initial_schema.sql, 002_...

```
1 -- product table
2 CREATE TABLE product (
3     id text PRIMARY KEY CHECK (ID ≠ '') NOT NULL,
4     name text NOT NULL CHECK (NAME ≠ ''),
5     description text NOT NULL,
6     price int NOT NULL CHECK (price ≥ 0),
7     created_at timestamp with time zone NOT NULL DEFAULT now(),
8     modified_at timestamp with time zone NOT NULL DEFAULT now()
9     -- If you want to use a soft delete strategy, you'll need something like:
10     -- deleted_at timestamp with time zone DEFAULT now()
11     -- or better: a product_history table to keep track of each change here.
12 );
13
14 COMMENT ON COLUMN product.id IS 'assume id is the barcode';
15 COMMENT ON COLUMN product.price IS 'price in the smaller subdivision possible (such as cents)';
16 CREATE INDEX product_name ON product(name text_pattern_ops);
17
18 — create above / drop below —
19
20 DROP TABLE product;
```

tern: standalone migration tool

<https://asciinema.org/a/450576>

Usage:
tern [command]

Available Commands:

code	Execute a code package command
help	Help about any command
init	Initialize a new tern project
migrate	Migrate the database
new	Generate a new migration
status	Print current migration status
version	Print version



Flags:
-h, --help help for tern

Use "tern [command] --help" for more information about a command.


```
henvic in ~/projects/gocode/src/github.com/henvic/pgxtutorial (main●) prod  
$ cd migrations
```

```
henvic in ~/projects/gocode/src/github.com/henvic/pgxtutorial/migrations (main●) prod  
$ tern m
```

Database layer

It might be useful to define an interface, even if bothersome.

What about mocks?



```
1 // DB layer.
2 //go:generate mockgen --build_flags=--mod=mod -package inventory -destination mock_db_test.go . DB
3 type DB interface {
4     // CreateProduct creates a new product.
5     CreateProduct(ctx context.Context, params CreateProductParams) error
6
7     // UpdateProduct updates an existing product.
8     UpdateProduct(ctx context.Context, params UpdateProductParams) error
9
10    // GetProduct returns a product.
11    GetProduct(ctx context.Context, id string) (*Product, error)
12
13    // ...
14
15    // TransactionContext returns a copy of the parent context which begins a transaction
16    // to PostgreSQL.
17    TransactionContext(ctx context.Context) (context.Context, error)
18
19    // Commit transaction from context.
20    Commit(ctx context.Context) error
21
22    // Rollback transaction from context.
23    Rollback(ctx context.Context) error
24
25    // WithAcquire returns a copy of the parent context which acquires a connection
26    // to PostgreSQL from pgxpool to make sure commands executed in series reuse the
27    // same database connection.
28    WithAcquire(ctx context.Context) (dbCtx context.Context, err error)
29
30    // Release PostgreSQL connection acquired by context back to the pool.
31    Release(ctx context.Context)
32 }
```

Testing strategy

Get superior returns from your tests:

- Favor writing integration test with a real implementation
- Use test doubles to test scenarios of database failure.

Real implementation	Test doubles (mocks, fakes...)
SELECT returning data or not found.	Database connection issue.
INSERT, UPDATE, etc.	Simulation of unexpected DB error.
Constraints checks.	Expensive operation (ask why first)

Inspiration: Software Engineering at Google: Lessons Learned from Programming Over Time

<https://amzn.to/3uqZWP7>



**Just like
Theranos'
blood tests...**

**Testing
interfaces for
the sake of it
is overrated.**

Some tools

- GoMock github.com/golang/mock
- go-cmp github.com/google/go-cmp
- scany github.com/georgysavva/scany
- pgtools github.com/hatch-studio/pgtools
- squirrel github.com/Masterminds/squirrel (I prefer plain SQL, though)

GoMock

github.com/golang/mock

```
1 ctrl := gomock.NewController(t)
2 m := inventory.NewMockDB(ctrl)
3 m.EXPECT().CreateProduct(gomock.Not(gomock.Nil()),
4     inventory.CreateProductParams{
5         ID: "simple",
6         Name: "product name",
7         Description: "product description",
8         Price: 150,
9     }).Return(errors.New("unexpected error"))
10 return m
```

Comparing values with go-cmp

github.com/google/go-cmp

- *A more powerful and safer alternative to reflect.DeepEqual for comparing whether two values are semantically equal.*

```
1 // Ignoring field generated automatically:
2 if !cmp.Equal(tt.want, got, cmpopts.IgnoreFields(inventory.ProductReview{}, "ID")) {
3     t.Errorf("value returned by Service.GetProductReview() doesn't match: %v", cmp.Diff(tt.want, got))
4 }
5
6 // Several ways to check for equality treating time values as special:
7 if !cmp.Equal(tt.want, got, cmpopts.EquateApproxTime(time.Minute)) {
8     t.Errorf("value returned by Service.GetProduct() doesn't match: %v", cmp.Diff(tt.want, got))
9 }
10
11 if !cmp.Equal(tt.want, got, cmpopts.IgnoreFields(inventory.Product{}, "CreatedAt", "ModifiedAt")) {
12     t.Errorf("value returned by DB.GetProduct() doesn't match: %v", cmp.Diff(tt.want, got))
13 }
14
15 if !cmp.Equal(tt.want, got, cmpopts.EquateApproxTime(time.Minute)) {
16     t.Errorf("value returned by Service.GetProduct() doesn't match: %v", cmp.Diff(tt.want, got))
17 }
18
19 if !cmp.Equal(tt.want, got, cmpopts.EquateApproxTime(time.Minute)) {
20     t.Errorf("value returned by Service.SearchProducts() doesn't match: %v", cmp.Diff(tt.want, got))
21 }
22
23 if !cmp.Equal(tt.want, got, cmpopts.IgnoreTypes(time.Time{})) {
24     t.Errorf("value returned by DB.GetProductReviews() doesn't match: %v", cmp.Diff(tt.want, got))
25 }
```

scany

```
package main

import (
    "context"

    "github.com/jackc/pgx/v4/pgxpool"

    "github.com/georgysavva/scany/pgxscan"
)

type User struct {
    ID    string
    Name  string
    Email string
    Age   int
}

func main() {
    ctx := context.Background()
    db, _ := pgxpool.Connect(ctx, "example-connection-url")

    var users []*User
    pgxscan.Select(ctx, db, &users, `SELECT id, name, email, age FROM users`)
    // users variable now contains data from all rows.
}
```


pgtools

github.com/hatch-studio/pgtools contains:

- Package to make creating SELECT for writing queries easier.
- Package sqltest, which makes writing integration tests very efficient.

Tests with sqltest

- Auto-discovery recognizes PostgreSQL configuration automatically.
- Creates a database with prefix test_ (safety) + normalized test names.
- Your tests can run in parallel just as fine.



sqltest + tern = easy tests <https://asciinema.org/a/450576>


```
1 var force = flag.Bool("force", false, "Force cleaning the database before starting")
2
3 func TestCreateProduct(t *testing.T) {
4     t.Parallel()
5     migration := sqltest.New(t, sqltest.Options{
6         Force: *force,
7         Path:  ".. / .. /migrations",
8     })
9     pool := migration.Setup(context.Background(), "")
10
11     db := &DB{
12         Postgres: pool,
13     }
14     // ...
15 }
```


Then you discover running its integration tests is hard.

```
$ go test ./...
```

```
-----  
/ --- FAIL: TestRollback (0.01s) \\  
| \\  
| sqltest.go:222: open postgres \\  
| connection: failed to connect to \\  
| `host=lol user=hatch \\  
| database=postgres`: hostname resolving \\  
| error (lookup lol: no such host) FAIL \\  
| FAIL presentation/internal/dao 0.193s \\  
\ FAIL /
```

```
 \  ^__^  
 \  (oo)\_____  
    (__)\\       )\\/\  
        ||----w |  
        ||     ||
```



\$ go test ./...

From zero to running running integration tests in few seconds. Try it:

1. Install or configure PostgreSQL [environment variables](#) so that psql works.
2. \$ export INTEGRATION_TESTDB=true
\$ go test -v ./...

Continuous Integration

Run your DB on any CI/CD system.

For example, with a pull request [workflow](#) on GitHub Actions, stateless.



```
1 name: Integration
2 on:
3   pull_request:
4     types: [opened, synchronize, reopened, ready_for_review]
5   push:
6     branches:
7       - main
8   permissions:
9     contents: read
10    pull-requests: read
11 jobs:
12   # Reference: https://docs.github.com/en/actions/guides/creating-postgresql-service-containers
13   postgres-test:
14     runs-on: ubuntu-latest
15     services:
16       postgres:
17         image: postgres
18         env:
19           POSTGRES_USER: runner
20           POSTGRES_PASSWORD: postgres
21           POSTGRES_DB: test_pgxtutorial
22         options: >-
23           --name postgres
24           --health-cmd pg_isready
25           --health-interval 10s
26           --health-timeout 5s
27           --health-retries 5
28         ports:
29           # Maps tcp port 5432 on service container to the host
30           - 5432:5432
31     env:
32       INTEGRATION_TESTDB: true
33       PGHOST: localhost
34       PGUSER: runner
35       PGPASSWORD: postgres
36       PGDATABASE: test_pgxtutorial
37     steps:
38       - uses: actions/checkout@v1
39       - uses: actions/setup-go@v1
40         with:
41           go-version: '1.17.x'
42       - name: Run Postgres tests
43         run: go test -v -race -count 1 -covermode atomic -coverprofile=profile.cov ./...
44       - name: Code coverage
45         if: ${github.event_name} != 'pull_request'
46         uses: shogo82148/actions-goveralls@v1
47         with:
48           path-to-profile: profile.cov
```

```
1 $ git clone https://github.com/henvic/pgxtutorial.git
```

```
2 $ cd pgxtutorial
```

```
1 # Run tests with:
```

```
2 $ INTEGRATION_TESTDB=true go test -v ./...
```

```
1 # To execute application, first create a database and run migrations.
```

```
2 $ psql -c "CREATE DATABASE pgxtutorial;"
```

```
3 $ export PGDATABASE=pgxtutorial
```

```
4 $ tern migrate -m ./migrations
```

```
5 # Then, run it (without installing):
```

```
6 $ go run ./cmd/pgxtutorial
```

```
7 2021/11/22 07:21:21 HTTP server listening at localhost:8080
```

```
8 2021/11/22 07:21:21 gRPC server listening at 127.0.0.1:8082
```

Proof it works <https://asciinema.org/a/450580>

```
api.Inventory@localhost:8082> call SearchProducts
```

```
query_string (TYPE_STRING) => chair
```

```
✓ min_price
```

```
min_price (TYPE_INT64) => 1
```

```
✓ max_price
```

```
max_price (TYPE_INT64) => 1000
```

```
✓ page
```

```
page (TYPE_INT32) => 1
```

```
{  
  "items": [  
    {  
      "description": "This is a nice chair.",  
      "id": "chair",  
      "name": "Nice office chair",  
      "price": "14"  
    }  
  ],  
  "total": 1  
}
```

```
api.Inventory@localhost:8082> █
```

Takeaways

- PostgreSQL is good
- There are many tools to help you
- Balance between integration and unit tests
- Continuous Integration is easier than you think

Thanks. Questions?

Henrique Vicente, HATCH Studio

- Article: <https://henvic.dev/posts/go-postgres/>
- Repository: <https://github.com/henvic/pgxtutorial>