

# gRPC

The load balancing gotchas and the  
connection pool solution



By: Ruba Zeibak (Stitch)


# Agenda

- What is gRPC? A short introduction
- HTTP/2 - The secret weapon of gRPC
- gRPC at Stitch
- The FLOW\_CONTROL error
- gRPC in k8s and load balancing
- Demo - it's GO time!

# RPC Architecture

## Remote Procedure Call

RPC protocols allow to invoke a function/method on a remote server and get the result in the same format regardless of where it is executed



# What is gRPC?

- gRPC is an open-source Remote Procedure Call (RPC) framework that is used for high-performance communication between services.
- It is an efficient way to connect services written in different languages with pluggable support for load balancing, tracing, health checking, and authentication



# Is g for Google?

GRPC Core: g stands for  
'g' stands for something different every gRPC  
release:  
[grpc.github.io](https://grpc.github.io)

- 1.0 'g' stands for 'gRPC'
- 1.1 'g' stands for 'good'
- 1.2 'g' stands for 'green'
- 1.3 'g' stands for 'gentle'
- 1.4 'g' stands for 'gregarious'
- 1.6 'g' stands for 'garcia'
- 1.7 'g' stands for 'gambit'
- 1.8 'g' stands for 'generous'
- 1.9 'g' stands for 'glossy'
- 1.10 'g' stands for 'glamorous'
- 1.11 'g' stands for 'gorgeous'
- 1.12 'g' stands for 'glorious'
- 1.13 'g' stands for 'gloriosa'
- 1.14 'g' stands for 'gladiolus'
- 1.15 'g' stands for 'glider'
- 1.16 'g' stands for 'gao'
- 1.17 'g' stands for 'gizmo'
- 1.18 'g' stands for 'goose'
- 1.19 'g' stands for 'gold'
- 1.20 'g' stands for 'godric'
- 1.21 'g' stands for 'gandalf'
- 1.22 'g' stands for 'gale'
- 1.23 'g' stands for 'gangnam'
- 1.24 'g' stands for 'ganges'
- 1.25 'g' stands for 'game'
- 1.26 'g' stands for 'gon'
- 1.27 'g' stands for 'guantao'
- 1.28 'g' stands for 'galactic'
- 1.29 'g' stands for 'gringotts'
- 1.30 'g' stands for 'gradius'
- 1.31 'g' stands for 'galore'
- 1.32 'g' stands for 'giggle'
- 1.33 'g' stands for 'geeky'
- 1.34 'g' stands for 'gauntlet'
- 1.35 'g' stands for 'gecko'
- 1.36 'g' stands for 'gummybear'
- 1.37 'g' stands for 'gilded'
- 1.38 'g' stands for 'guadalupe\_river\_park\_conservancy'
- 1.39 'g' stands for 'goofy'
- 1.40 'g' stands for 'guileless'

# Protocol Buffers

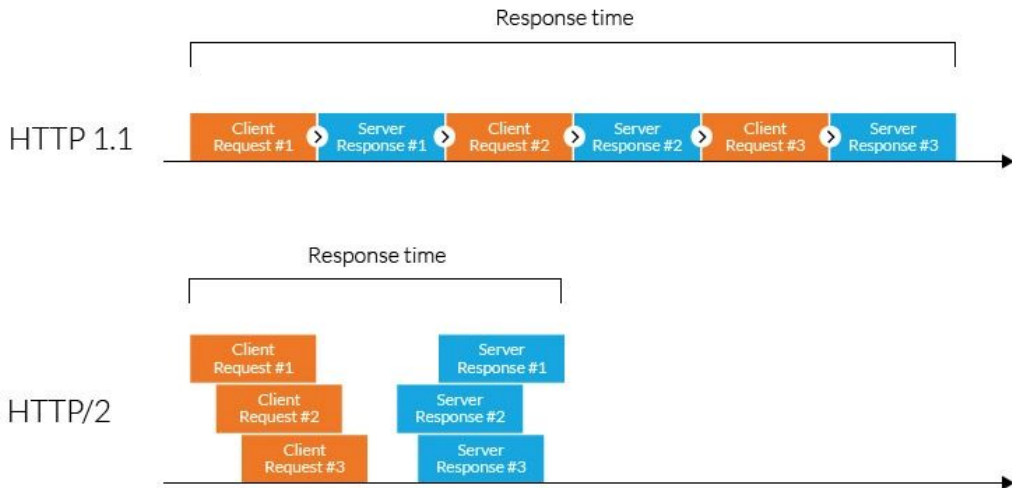
gRPC uses protocol buffers as the interface definition language for serialization and communication instead of JSON/XML.




# HTTP/2 - The secret weapon of gRPC

## HTTP/1.1 vs. HTTP/2 Protocol

- Binary protocols
- Multiplexing
- Header compression
- Server push
- Increased security



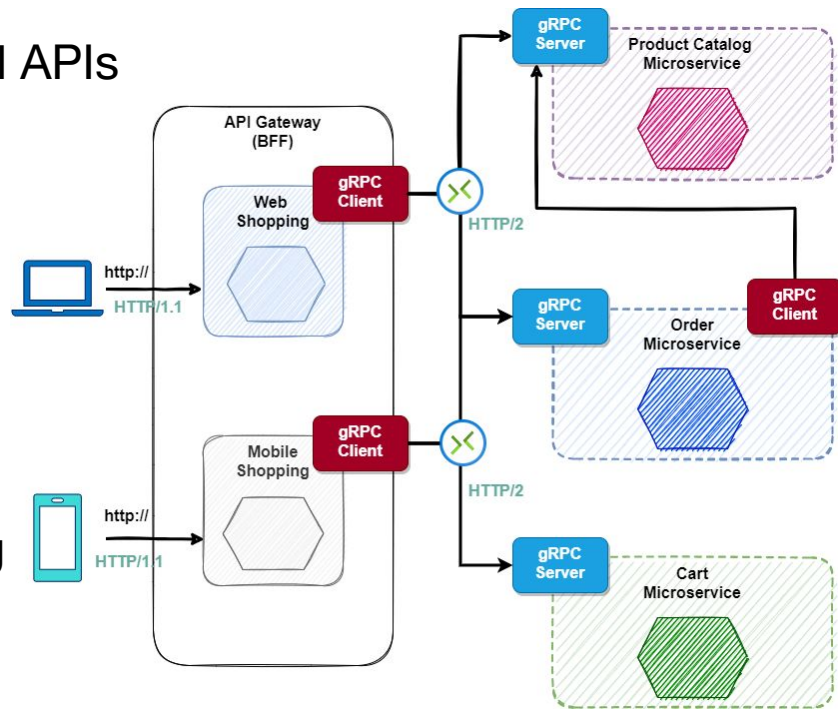
# Types of gRPC



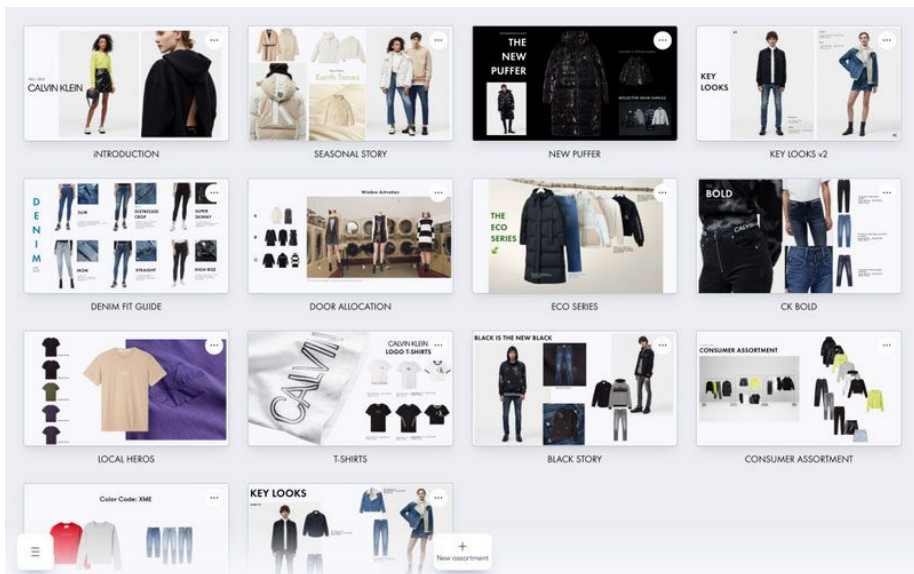


# What is gRPC good for?

- Faster compared to JSON based RESTful APIs
- Strong type and well-defined interface
- Polyglot
- Stream support
- Built in support for auth, load balancing, encryption, compression and error handling



# The Digital Showroom at Stitch




## DENIM

Fall 2022



# gRPC at Stitch



# HTTP/2 Flow Control

## The unexpected error

Flow Control allows a receiver to stop a sender from sending it data if it isn't yet ready to process, perhaps because it's too busy to process any more incoming data.

**level=error msg="finished unary call with code Internal" api=grpc error="rpc error: code = Internal desc = create the media: stream error: stream ID 7977; FLOW\_CONTROL\_ERROR; received from peer" grpc.code=Internal**

# gRPC and k8s

isn't that straight forward after all

<https://kubernetes.io/blog/2018/11/07/grpc-load-balancing-on-kubernetes-without-tears/>

---

## gRPC Load Balancing on Kubernetes without Tears

Wednesday, November 07, 2018

**Author:** William Morgan (Buoyant)

# Load balancing gRPC

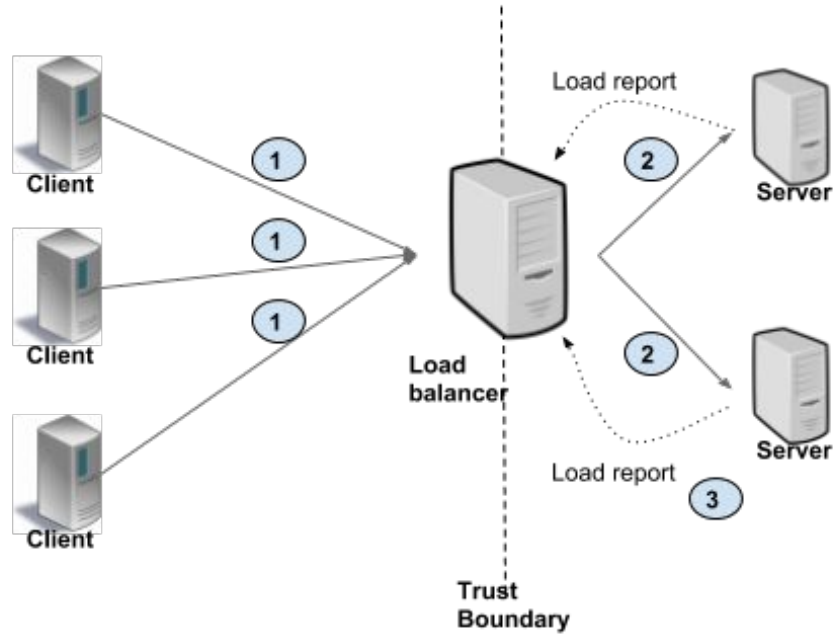
## In k8s

In order to do gRPC load balancing, we need to shift from connection balancing (L3/L4 which is the k8s default) to *request* balancing (L5/L7).

In other words, we need to open an HTTP/2 connection to each destination, and balance *requests* across these connections.

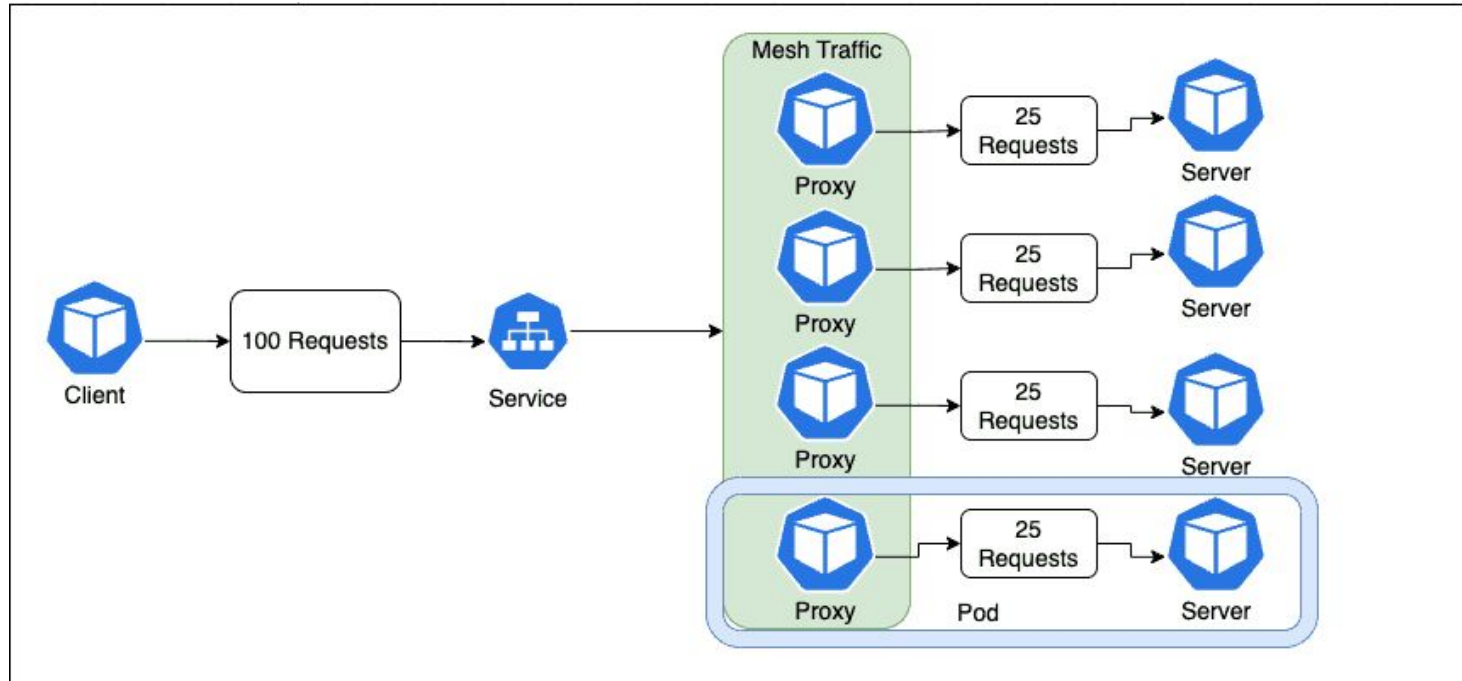
# Proxy Load balancing

Also known as **Server-Side LB**



# Proxy LB in K8s


Using service mesh like Linkerd or istio






# Client Side Load balancing

## Thick LB



## Lookaside LB



# Client side LB in K8s

Client-side load balancing using headless service

You can do client-side **round-robin** load-balancing using **Kubernetes headless service**. This simple load balancing works out of the box with gRPC.

The downside is that it does not take into account the load on the server.

# The gRPC connection pool solution

**The chosen solution:** A gRPC connection pool on client side combined with a server side TCP (Layer4) load balancer.

This will create a pool of client connections initially, and re-use this pool of connections for subsequent gRPC requests.

Using <https://github.com/processout/grpc-go-pool>

# It's finally Go time



```
syntax = "proto3";
```

```
package media;
```

```
message DownloadRequest {
```

```
    string url = 1;
```

```
}
```

```
message MediaDimensions {
```

```
    //height of the image
```

```
    double height = 1;
```

```
    //width of the image
```

```
    double width = 2;
```

```
}
```

```
service Medias {
```

```
    //Download
```

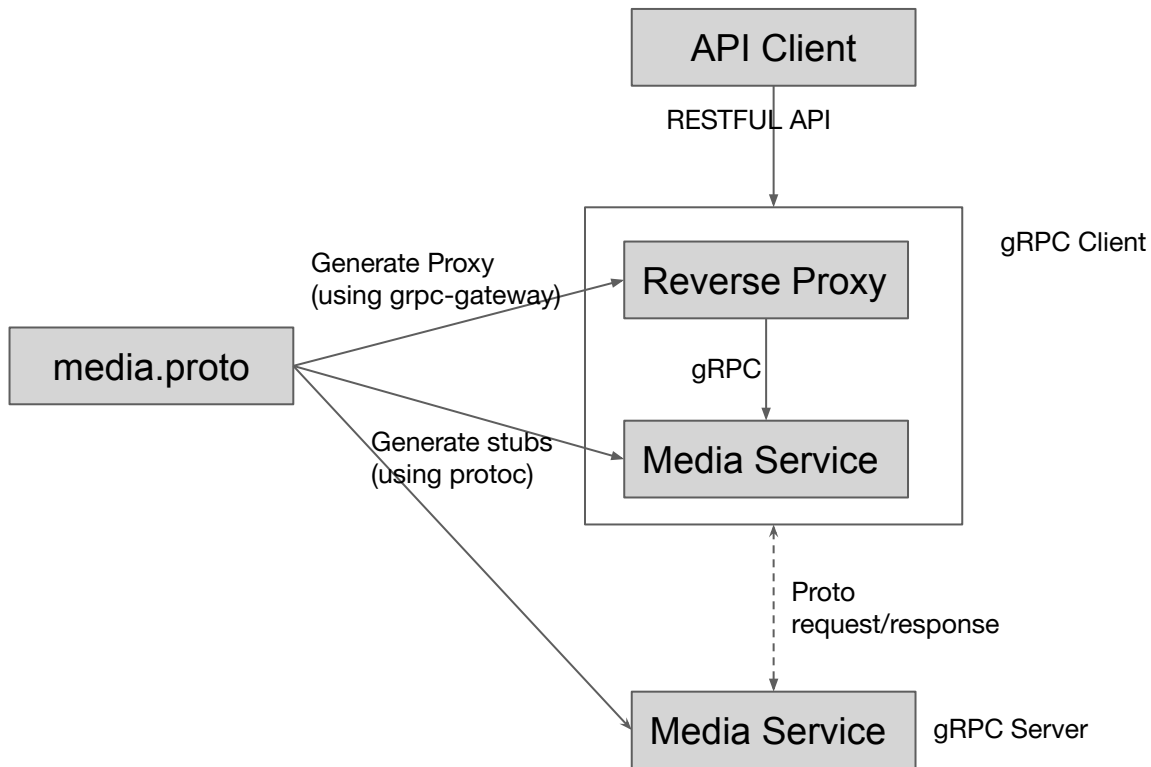
```
    //
```

```
    // Download Media example for Go meetup Amsterdam May 2022
```

```
    rpc Download (DownloadRequest) returns (MediaDimensions);
```

```
}
```

# gRPC at Stitch



```

// RegisterMediasHandlerClient registers the http handlers for service Medias
// to "mux". The handlers forward requests to the grpc endpoint over the given implementation of "MediasClient".
// Note: the gRPC framework executes interceptors within the gRPC handler. If the passed in "MediasClient"
// doesn't go through the normal gRPC flow (creating a gRPC client etc.) then it will be up to the passed in
// "MediasClient" to call the correct interceptors.

func RegisterMediasHandlerClient(ctx context.Context, mux *runtime.ServeMux, client MediasClient) error {
    mux.Handle(meth: "POST", pattern_Medias_Download_0, func(w http.ResponseWriter, req *http.Request,
        pathParams map[string]string) {
        ctx, cancel := context.WithCancel(req.Context())
        defer cancel()
        inboundMarshaler, outboundMarshaler := runtime.MarshalerForRequest(mux, req)
        rctx, err := runtime.AnnotateContext(ctx, mux, req, rpcMethodName: "/media.Medias/Download")
        if err != nil {
            runtime.HTTPError(ctx, mux, outboundMarshaler, w, req, err)
            return
        }

        resp, md, err := request_Medias_Download_0(rctx, inboundMarshaler, client, req, pathParams)
        ctx = runtime.NewServerMetadataContext(ctx, md)
        if err != nil {
            runtime.HTTPError(ctx, mux, outboundMarshaler, w, req, err)
            return
        }

        forward_Medias_Download_0(ctx, mux, outboundMarshaler, w, req, resp, mux.GetForwardResponseOptions()...)

    })

    return nil
}

```

```
// MediasClient is the client API for Medias service.
//
// For semantics around ctx use and closing/ending streaming RPCs, please refer to https://pkg.go.dev/google.golang.org/grpc/?tab=doc#ClientConn.NewStream.
```

```
type MediasClient interface {
    //Download
    //
    // Download Media example for Go meetup Amsterdam May 2022
    Download(ctx context.Context, in *DownloadRequest, opts ...grpc.CallOption) (*MediaDimensions, error)
}
```

```
type mediasClient struct {
    cc grpc.ClientConnInterface
}
```

```
func NewMediasClient(cc grpc.ClientConnInterface) MediasClient {
    return &mediasClient{ cc: cc}
```

```
func (c *mediasClient) Download(ctx context.Context, in *DownloadRequest, opts ...grpc.CallOption) (*MediaDimensions, error) {
    out := new(MediaDimensions)
    err := c.cc.Invoke(ctx, method: "/media.Medias/Download", in, out, opts...)
    if err != nil : nil, err
    return out, nil
}
```



```
func NewMediasClient(cc grpc.ClientConnInterface) MediasClient {  
    return &mediasClient{ cc: cc}  
}
```

```
// ClientConnInterface defines the functions clients need to perform unary and  
// streaming RPCs. It is implemented by *ClientConn, and is only intended to  
// be referenced by generated code.  
type ClientConnInterface interface {  
    // Invoke performs a unary RPC and returns after the response is received  
    // into reply.  
    Invoke(ctx context.Context, method string, args interface{}, reply interface{}, opts ...CallOption) error  
    // NewStream begins a streaming RPC.  
    NewStream(ctx context.Context, desc *StreamDesc, method string, opts ...CallOption) (ClientStream, error)  
}
```

```

//implements grpc.ClientConnInterface
type connClientWithPool struct {
    endpoint string
    pool      *grpcpool.Pool
}

// Invoke performs a unary RPC and returns after the response is received
// into reply.
func (cp connClientWithPool) Invoke(ctx context.Context, method string, args interface{}, reply interface{},
    opts ...grpc.CallOption) error {
    conn, err := cp.pool.Get(ctx)
    if err != nil {
        err = errors.Wrapf(err, format: "get pool connection method %v, on endpoint %v on invoke", method, cp.endpoint)
        return err
    }

    defer func() {
        err := conn.Close()
        if err != nil {
            logr.Warn(errors.Wrap(err, message: "return connection back to pool"))
        }
    }()

    logr.Debug( args...: "invoking the method")
    if err := conn.Invoke(ctx, method, args, reply, opts...); err != nil {
        err = errors.Wrapf(err, format: "invoke method %v, on endpoint %v", method, cp.endpoint)
        logr.Warn(err)
        return err
    }

    return nil
}

```

```
func main() {  
    ctx := context.Background()  
  
    cfg := config{}  
    if err := envconfig.Process( prefix: "", &cfg); err != nil {  
        log.Fatal(errors.Wrap(err, message: "get env configs"))  
    }  
  
    mux := runtime.NewServeMux()  
    opts := []grpc.DialOption{  
        grpc.WithConnectParams(grpc.ConnectParams{  
            Backoff: backoff.DefaultConfig,  
        })),  
    }  
  
    if err := initMedia(ctx, mux, cfg, opts); err != nil {  
        log.Fatal(errors.Wrap(err, message: "init endpoints"))  
    }  
  
    server := pkghttp.NewServer(logr, mux, basepath: "/", cfg.RESTAPIHost, cfg.RESTAPIPort)  
    if err := server.Close(); err != nil {  
        logr.Error(err)  
    }  
}
```

```
func initMedia(ctx context.Context, mux *runtime.ServeMux, cfg config, opts []grpc.DialOption) error {
    poolCfg := poolConfig{
        max:          cfg.GRPCPoolMax,
        init:         cfg.GRPCPoolInit,
        timeoutInSeconds: cfg.GRPCPoolIdleTimeoutSeconds,
    }
    endpoint := net.JoinHostPort(cfg.MediaHost, cfg.MediaPort)
    conn, err := grpcDialWithPool(endpoint, poolCfg, opts)
    if err != nil {
        return errors.Wrap(err, message: "connect to media")
    }
    // Methods below always return nil error
    _ = media.RegisterMediasHandlerClient(ctx, mux, media.NewMediasClient(conn))
    return nil
}
```

```
// Return a new grpc client connection that implements grpc.ClientConnInterface with connection pool that returns the
// first available client connection.

func grpcDialWithPool(endpoint string, poolCfg poolConfig, opts []grpc.DialOption) (conn *connClientWithPool, err error) {
    pool, ok := pools[endpoint]
    if !ok {
        pool, err = grpcpool.New(factory(endpoint, opts), poolCfg.init, poolCfg.max, poolCfg.timeoutInSeconds)
        if nil != err {
            logr.Errorln(err)
            return conn: nil, err
        }
        pools[endpoint] = pool
    }

    logr.Infof("format: \"grpcpool endpoint:%s, capacity:%d,available:%d\", endpoint, pool.Capacity(), pool.Available())

    return &connClientWithPool{
        endpoint: endpoint,
        pool:     pool,
    }, err
}
```

```
func factory(endpoint string, opts []grpc.DialOption) func() (*grpc.ClientConn, error) {  
    return func() (*grpc.ClientConn, error) {  
        ctx, cancel := context.WithTimeout(context.Background(), 10*time.Second)  
        defer cancel()  
        conn, err := grpc.DialContext(ctx, endpoint, opts...)  
        if err != nil {  
            return nil, err  
        }  
        return conn, nil  
    }  
}
```

# Was this a good choice?

## **Advantages:**

- Reusing of existing L4 K8s LB - and hence all pods share the load
- Still uses Multiplexing since the `pool.Get()` returns the next available client which reuses all the connections in the pool
- No more `FLOW_CONTROL` errors

## **Disadvantages:**

- Thickening the client logic
- Not as scalable as a proper LB solution

# Questions?

